

(Assignment-01)

Numpy

```
from genericpath import getsize
import sys as sys
a="Hello, man how are you "
sys.getsizeof(a)
fact = 1
for i in range(1,11):
    fact = fact*i
print(sys,a,fact)
```

```
<module 'sys' (built-in)> Hello, man how are you 1
<module 'sys' (built-in)> Hello, man how are you 2
<module 'sys' (built-in)> Hello, man how are you 6
<module 'sys' (built-in)> Hello, man how are you 24
<module 'sys' (built-in)> Hello, man how are you 120
<module 'sys' (built-in)> Hello, man how are you 720
<module 'sys' (built-in)> Hello, man how are you 5040
<module 'sys' (built-in)> Hello, man how are you 40320
<module 'sys' (built-in)> Hello, man how are you 362880
<module 'sys' (built-in)> Hello, man how are you 3628800
```

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
x = []
print(type(x))
```

```
x = ["hello", "kaho", "chalo", "suno", 67]
# print(x[1:4])
# print(len(x))
x[4] = 'hello habibi'
# print(x)
x.append('hey mister')
# print(x)
x[4] = ["hey", 5, "say hi"]
```

```
x[4][0] = [2,6,7,[4,5]]
print(x[:5:1])
for t in x:
    print(t,end=" ")
```

```
<class 'list'>
['hello', 'kaho', 'chalo', 'suno', [[2, 6, 7, [4, 5]], 5, 'say hi']]
hello kaho chalo suno [[2, 6, 7, [4, 5]], 5, 'say hi'] hey mister
```

```
import numpy as np
list1 = [1,2,3,4,5,6,7,78,8,8,6,5,4,53,[1,5]]
list2 = []
print(dir(list1))
print(len(dir(list1)))
print(list1)
x1 = np.array(list1)
y = x
print(y)
print(x1)
for i in list1:
    print(i)
print(list2)
list2.append(list1)
print(list2)
['__add__', '__class__', '__contains__', '__delattr__', '__delitem__', '__dir__', '__doc__', '__eq__', '__format__', '__ge__',
46
[1, 2, 3, 4, 5, 6, 7, 78, 8, 8, 6, 5, 4, 53, [1, 5]]
[[54, 5], 6, 7]
[1 2 3 4 5 6 7 78 8 8 6 5 4 53 list([1, 5])]
```

```
import numpy as np
list1 = [1,2,3,25,4,53,[1,5]]
list2 = []
print(dir(list1))
print(len(dir(list1)))
print(list1)
x1 = np.array(list1)
y = x
print(y)
print(x1)
```

```

for i in list1:
    print(i)
print(list2)
list2.append(list1)
print(list2)

import numpy as np
list1 = [1,2,3,25,4,53,[1,5]]
list2 = []
print(dir(list1))
print(len(dir(list1)))
print(list1)
x1 = np.array(list1)
y = x
print(y)
print(x1)
for i in list1:
    print(i)
print(list2)
list2.append(list1)
print(list2)


['_add_', '__class__', '__class_getitem__', '__contains__', '__delattr__', '__delitem__', '__dir__', '__doc__', '__eq__', '__format__', '__ge_'
47
[1, 2, 3, 25, 4, 53, [1, 5]]
['hello', 'kaho', 'chalo', 'suno', [[2, 6, 7, [4, 5]], 5, 'say hi'], 'hey mister']
[1 2 3 25 4 53 list([1, 5])]
1
2
3
25
4
53
[1, 5]
[]
[[1, 2, 3, 25, 4, 53, [1, 5]]]
<ipython-input-9-381b22bfe96a>:7: VisibleDeprecationWarning: Creating an ndarray from ragged nested sequences (which is a list-or-tuple of lists-c
x1 = np.array(list1)

```

(Assignment-02) Matplotlib

```
import matplotlib.pyplot as plt
import numpy as np
x = np.array([4,15])
y = np.array([10,150])
plt.plot(x,y)
plt.show()
```


Output :-



```
import matplotlib.pyplot as plt
import numpy as np
x = np.linspace(-2,2,20)
print(x)
y = 4*x**2
plt.plot(x,y)
plt.show()
```


output:-

```
[-2.          -1.78947368 -1.57894737 -1.36842105 -1.15789474 -0.94736842
 -0.73684211 -0.52631579 -0.31578947 -0.10526316  0.10526316  0.31578947
 0.52631579  0.73684211  0.94736842  1.15789474  1.36842105  1.57894737
 1.78947368  2.         ]
```




```
x = np.array([2,8])
y = np.array([4,15])
plt.plot(x,y,'o')
plt.show()
```

Output :-




```
x = np.array([1,3,5,10,4,12])
y = np.array([0,4,5,6,5,30])
plt.plot(x,y,marker='o')
plt.show()
```

output:-



```
y = np.array([2,3,1,0,12,23,11])
plt.plot(y)
plt.show()
```


output:-



```
y = np.array([2,3,1,0,12,23,11])
```


```
plt.plot(y,marker='o')
plt.show()
```

output:-




```
y = np.array([2,6,1,10,12,0])
plt.plot(y,marker='H')
plt.show()
```

Output:-




```
import matplotlib.pyplot as plt
import numpy as np
ypoints = np.array([1,4,1,4,1])
plt.plot(ypoints, 'd:r')
plt.show()
```




```
y = np.array([1,4,2,5,4,2,5])
plt.plot(y, marker='o', ms=30, mec = 'y')
plt.show()
```

output:-



```
y = np.array([1,4,2,5,6,4,2,5,10])
line, = plt.plot(y, marker='o', ms=30, mfc = 'r')
line.set_color("green")
```



(Assignment-03) Confusion Matrix


```
from sklearn.datasets import load_iris
import sklearn;
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
iris = load_iris();
x = iris.data
y = iris.target
x_train, x_test, y_train, y_test = train_test_split(x,y)
dtc = LinearRegression()
dtc.fit(x_train, y_train)
print(y_train)
print(y_test)
```

output:-

```
[1 0 2 1 1 1 2 0 1 0 2 0 1 1 1 1 0 0 1 1 1 2 2 0 0 0 2 2 1 2 2 1 0 0 0 1 1
 1 2 2 0 2 0 1 2 0 1 0 1 0 1 2 0 1 1 2 1 2 0 2 1 1 0 2 0 2 2 2 0 1 2 0 0 0
 0 1 1 0 0 0 0 2 0 2 2 1 2 0 0 1 2 1 2 1 2 0 1 1 0 0 0 2 0 0 2 1 2 2 2 0 2
 2]
[2 1 2 1 0 0 0 0 2 0 1 2 1 1 1 2 1 1 2 0 1 0 2 2 0 2 1 1 1 2 0 2 1 2 2 1 2
 0]
```

```
import matplotlib.pyplot as plt
import numpy
from sklearn import metrics
actual = numpy.random.binomial(1,.9,size = 1000)
predicted = numpy.random.binomial(1,.9,size = 1000)
confusion_matrix = metrics.confusion_matrix(actual, predicted)
cm_display = metrics.ConfusionMatrixDisplay(confusion_matrix =
confusion_matrix, display_labels = [False, True])
cm_display.plot()
plt.show()
```

output:-



(Assignment-04)

Simple Linear Regression

```
import matplotlib.pyplot as plt
from scipy import stats


x = [5,7,8,7,2,17,2,9,4,11,12,9,6]
y = [99,86,87,88,111,86,103,87,94,78,77,85,86]

slope, intercept, r, p, std_err = stats.linregress(x, y)

def myfunc(x):
    return slope * x + intercept

mymodel = list(map(myfunc, x))

plt.scatter(x, y)
plt.plot(x, mymodel)
plt.show()
```



```

# importing the dataset
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

dataset = pd.read_csv('Salary_Data.csv')
dataset.head()

# data preprocessing
X = dataset.iloc[:, :-1].values #independent variable array
y = dataset.iloc[:, 1].values #dependent variable vector

# splitting the dataset
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test =
train_test_split(X, y, test_size=1/3, random_state=0)

# fitting the regression model
from sklearn.linear_model import LinearRegression
regressor = LinearRegression()
regressor.fit(X_train, y_train) #actually produces the linear eqn for the da

# predicting the test set results
y_pred = regressor.predict(X_test)
y_pred

y_test

# visualizing the results
#plot for the TRAIN

plt.scatter(X_train, y_train, color='red') # plotting the observation line
plt.plot(X_train, regressor.predict(X_train), color='blue') # plotting the
regression line
plt.title("Salary vs Experience (Training set)") # stating the title of the
graph

plt.xlabel("Years of experience") # adding the name of x-axis
plt.ylabel("Salaries") # adding the name of y-axis
plt.show() # specifies end of graph

#plot for the TEST

plt.scatter(X_test, y_test, color='red')
plt.plot(X_train, regressor.predict(X_train), color='blue') # plotting the

```

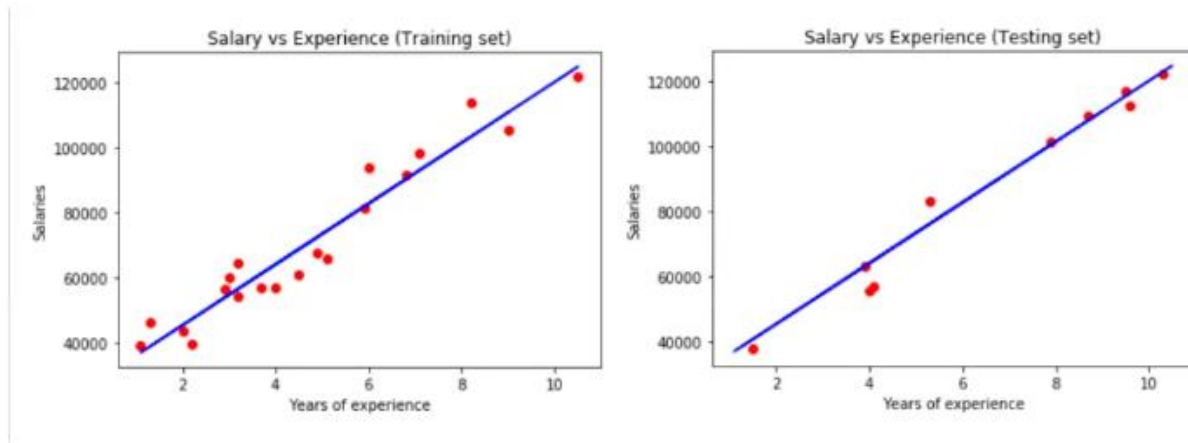
```

regression line
plt.title("Salary vs Experience (Testing set)")

plt.xlabel("Years of experience")
plt.ylabel("Salaries")
plt.show()

```

Output:-



Multiple Linear Regression

```

import pandas as pd
from sklearn import linear_model
import statsmodels.api as sm

data = {'year': [2017,2017,2017,2017,2017,2017,2017,2017,2017,2017,20
17,2017,2016,2016,2016,2016,2016,2016,2016,2016,2016,2016,2016,20
16,2016],
        'month': [12,11,10,9,8,7,6,5,4,3,2,1,12,11,10,9,8,7,6,5,
4,3,2,1],
        'interest_rate': [2.75,2.5,2.5,2.5,2.5,2.5,2.5,2.5,2.25,2.25
,2.25,2,2,2,1.75,1.75,1.75,1.75,1.75,1.75,1.75,1.75,1.75,1.75,1.
75],
        'unemployment_rate': [5.3,5.3,5.3,5.3,5.4,5.6,5.5,5.5,5.
5,5.6,5.7,5.9,6,5.9,5.8,6.1,6.2,6.1,6.1,6.1,5.9,6.2,6.2,6.1],
        'index_price': [1464,1394,1357,1293,1256,1254,1234,1195,
1159,1167,1130,1075,1047,965,943,958,971,949,884,866,876,822,704
,719]
        }

df = pd.DataFrame(data)

x = df[['interest_rate','unemployment_rate']]
y = df['index_price']

```

```

# with sklearn
regr = linear_model.LinearRegression()
regr.fit(x, y)

print('Intercept: \n', regr.intercept_)
print('Coefficients: \n', regr.coef_)

# with statsmodels
x = sm.add_constant(x) # adding a constant

model = sm.OLS(y, x).fit()
predictions = model.predict(x)

print_model = model.summary()
print(print_model)

```

output :-

```

Intercept:
1798.4039776258544

```

```

Coefficients:
[ 345.54008701 -250.14657137]

```

OLS Regression Results

```

=====
Dep. Variable:          index_price    R-squared:                0.898
Model:                  OLS          Adj. R-squared:            0.888
Method:                 Least Squares  F-statistic:              92.07
Date:                  Mon, 10 Apr 2023  Prob (F-statistic):      4.04e-11
Time:                  16:52:30       Log-Likelihood:           -134.61
No. Observations:      24            AIC:                     275.2
Df Residuals:          21            BIC:                     278.8
Df Model:               2
Covariance Type:       nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	1798.4040	899.248	2.000	0.059	-71.685	3668.493
interest_rate	345.5401	111.367	3.103	0.005	113.940	577.140
unemployment_rate	-250.1466	117.950	-2.121	0.046	-495.437	-4.856

```

=====
Omnibus:                 2.691    Durbin-Watson:              0.530
Prob(Omnibus):           0.260    Jarque-Bera (JB):         1.551
Skew:                   -0.612    Prob(JB):                 0.461
Kurtosis:                3.226    Cond. No.                 394.
=====

```

```

.. .

```

(Assignment-05) Logistic Regression

```
import numpy
from sklearn import linear_model

X = numpy.array([3.78, 2.44, 2.09, 0.14, 1.72, 1.65, 4.92, 4.37,
4.96, 4.52, 3.69, 5.88]).reshape(-1,1)
y = numpy.array([0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1, 1])

logr = linear_model.LogisticRegression()
logr.fit(X,y)

def logit2prob(logr, X):
    log_odds = logr.coef_ * X + logr.intercept_
    odds = numpy.exp(log_odds)
    probability = odds / (1 + odds)
    return(probability)

print(logit2prob(logr, X))
output:-
[[0.60749955]
 [0.19268876]
 [0.12775886]
 [0.00955221]
 [0.08038616]
 [0.07345637]
 [0.88362743]
 [0.77901378]
 [0.88924409]
 [0.81293497]
 [0.57719129]
 [0.96664243]]

import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn import datasets
iris = datasets.load_iris()
X = iris.data[:, :2]
y = (iris.target != 0) * 1
plt.figure(figsize=(6, 6))
plt.scatter(X[y == 0][:, 0], X[y == 0][:, 1], color='g', label='0')
plt.scatter(X[y == 1][:, 0], X[y == 1][:, 1], color='y', label='1')
plt.legend();
class LogisticRegression:
    def __init__(self, lr=0.01, num_iter=100000, fit_intercept=True, verbos
e=False):
```

```

        self.lr = lr
        self.num_iter = num_iter
        self.fit_intercept = fit_intercept
        self.verbose = verbose
    def __add_intercept(self, X):
        intercept = np.ones((X.shape[0], 1))
        return np.concatenate((intercept, X), axis=1)
    def __sigmoid(self, z):
        return 1 / (1 + np.exp(-z))
    def __loss(self, h, y):
        return (-y * np.log(h) - (1 - y) * np.log(1 - h)).mean()
    def fit(self, X, y):
        if self.fit_intercept:
            X = self.__add_intercept(X)

        self.theta = np.zeros(X.shape[1])
        for i in range(self.num_iter):
            z = np.dot(X, self.theta)
            h = self.__sigmoid(z)
            gradient = np.dot(X.T, (h - y)) / y.size
            self.theta -= self.lr * gradient
            z = np.dot(X, self.theta)
            h = self.__sigmoid(z)
            loss = self.__loss(h, y)
            if (self.verbose == True and i % 10000 == 0):
                print(f'loss: {loss} \t')


        def predict_prob(self, X)
    if self.fit_intercept:
        X = self.__add_intercept(X)
        return self.__sigmoid(np.dot(X, self.theta))
    def predict(self, X):
        return self.predict_prob(X).round()

    model = LogisticRegression(lr=0.1, num_iter=300000)
    preds = model.predict(X)
    (preds == y).mean()

    plt.figure(figsize=(10, 6))
    plt.scatter(X[y == 0][:, 0], X[y == 0][:, 1], color='g', label='0')
    plt.scatter(X[y == 1][:, 0], X[y == 1][:, 1], color='y', label='1')
    plt.legend()
    x1_min, x1_max = X[:, 0].min(), X[:, 0].max(),
    x2_min, x2_max = X[:, 1].min(), X[:, 1].max(),

```

```
xx1, xx2 = np.meshgrid(np.linspace(x1_min, x1_max), np.linspace(x2_min, x2_max))
grid = np.c_[xx1.ravel(), xx2.ravel()]
probs = model.predict_prob(grid).reshape(xx1.shape)
plt.contour(xx1, xx2, probs, [0.5], linewidths=1, colors='red');
```



```
import sklearn
from sklearn import datasets
from sklearn import linear_model
from sklearn import metrics
from sklearn.model_selection import train_test_split
digits = datasets.load_digits()
X = digits.data
y = digits.target
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.4, random_state=1)
digreg = linear_model.LogisticRegression()
digreg.fit(X_train, y_train)
y_pred = digreg.predict(X_test)
print("Accuracy of Logistic Regression model is:",
metrics.accuracy_score(y_test, y_pred)*100)
```

Output

```
Accuracy of Logistic Regression model is: 95.6884561891516
```

(Assignment-06) Desition Tree

```
# Importing the required packages
import numpy as np
import pandas as pd
from sklearn.metrics import confusion_matrix
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score
from sklearn.metrics import classification_report

# Function importing Dataset
def importdata():
    balance_data = pd.read_csv(
        'https://archive.ics.uci.edu/ml/machine-learning-'+
        'databases/balance-scale/balance-scale.data',
        sep= ',', header = None)

    # Printing the dataswet shape
    print ("Dataset Length: ", len(balance_data))
    print ("Dataset Shape: ", balance_data.shape)

    # Printing the dataset obseravtions
    print ("Dataset: ",balance_data.head())
    return balance_data

# Function to split the dataset
def splitdataset(balance_data):

    # Separating the target variable
    X = balance_data.values[:, 1:5]
    Y = balance_data.values[:, 0]

    # Splitting the dataset into train and test
    X_train, X_test, y_train, y_test = train_test_split(
        X, Y, test_size = 0.3, random_state = 100)

    return X, Y, X_train, X_test, y_train, y_test

# Function to perform training with giniIndex.
def train_using_gini(X_train, X_test, y_train):

    # Creating the classifier object
    clf_gini = DecisionTreeClassifier(criterion = "gini",
        random_state = 100,max_depth=3, min_samples_leaf=5)

    # Performing training
```



```

clf_gini.fit(X_train, y_train)
return clf_gini

# Function to perform training with entropy.
def tarin_using_entropy(X_train, X_test, y_train):

    # Decision tree with entropy
    clf_entropy = DecisionTreeClassifier(
        criterion = "entropy", random_state = 100,
        max_depth = 3, min_samples_leaf = 5)

    # Performing training
    clf_entropy.fit(X_train, y_train)
    return clf_entropy

# Function to make predictions
def prediction(X_test, clf_object):

    # Prediction on test with giniIndex
    y_pred = clf_object.predict(X_test)
    print("Predicted values:")
    print(y_pred)
    return y_pred

# Function to calculate accuracy
def cal_accuracy(y_test, y_pred):

    print("Confusion Matrix: ",
          confusion_matrix(y_test, y_pred))

    print ("Accuracy : ",
           accuracy_score(y_test,y_pred)*100)

    print("Report : ",
          classification_report(y_test, y_pred))

# Driver code
def main():

    # Building Phase
    data = importdata()
    X, Y, X_train, X_test, y_train, y_test = splitdataset(data)
    clf_gini = train_using_gini(X_train, X_test, y_train)
    clf_entropy = tarin_using_entropy(X_train, X_test, y_train)

    # Operational Phase
    # print("Results Using Gini Index:")

```

```

# Prediction using gini
# y_pred_gini = prediction(X_test, clf_gini)
# cal_accuracy(y_test, y_pred_gini)

print("Results Using Entropy:")
# Prediction using entropy
y_pred_entropy = prediction(X_test, clf_entropy)
cal_accuracy(y_test, y_pred_entropy)

# Calling main function
if __name__=="__main__":
    main()

```

Output :-

```

Dataset Length: 625
Dataset Shape: (625, 5)
Dataset:   0  1  2  3  4
0  B  1  1  1  1
1  R  1  1  1  2
2  R  1  1  1  3
3  R  1  1  1  4
4  R  1  1  1  5
Results Using Entropy:
Predicted values:
['R' 'L' 'R' 'L' 'R' 'L' 'R' 'L' 'R' 'R' 'R' 'R' 'L' 'L' 'R' 'L' 'R' 'L'
 'L' 'R' 'L' 'R' 'L' 'L' 'R' 'L' 'R' 'L' 'R' 'L' 'R' 'L' 'R' 'L' 'L' 'L'
 'L' 'L' 'R' 'L' 'R' 'L' 'R' 'L' 'R' 'R' 'L' 'L' 'R' 'L' 'L' 'R' 'L' 'L'
 'R' 'L' 'R' 'R' 'L' 'R' 'R' 'R' 'L' 'L' 'R' 'L' 'L' 'R' 'L' 'L' 'L' 'R'
 'R' 'L' 'R' 'L' 'R' 'R' 'R' 'L' 'R' 'L' 'L' 'L' 'L' 'R' 'R' 'L' 'R' 'L'
 'R' 'R' 'L' 'L' 'L' 'R' 'R' 'L' 'L' 'L' 'R' 'L' 'L' 'R' 'R' 'R' 'R' 'R'
 'R' 'L' 'R' 'L' 'R' 'R' 'L' 'R' 'R' 'L' 'R' 'R' 'L' 'R' 'R' 'R' 'L' 'L'
 'L' 'L' 'L' 'R' 'R' 'R' 'R' 'L' 'R' 'R' 'R' 'L' 'L' 'R' 'L' 'R' 'L' 'R'
 'L' 'R' 'R' 'L' 'L' 'R' 'L' 'R' 'R' 'R' 'R' 'R' 'L' 'R' 'R' 'R' 'R' 'R'
 'R' 'L' 'R' 'L' 'R' 'R' 'L' 'R' 'L' 'R' 'L' 'R' 'L' 'L' 'L' 'L' 'L' 'R'
 'R' 'R' 'L' 'L' 'L' 'R' 'R' 'R']
Confusion Matrix: [[ 0  6  7]
 [ 0 63 22]
 [ 0 20 70]]
Accuracy : 70.74468085106383
Report :

```


	precision	recall	f1-score	support
B	0.00	0.00	0.00	13
L	0.71	0.74	0.72	85
R	0.71	0.78	0.74	90
accuracy		0.71		188
macro avg	0.47	0.51	0.49	188
weighted avg	0.66	0.71	0.68	188

(Assignment-07) KNN Implementation

```
import numpy as nm;
import matplotlib.pyplot as plt;
import pandas as pd
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import train_test_split
from sklearn.datasets import load_iris
import numpy as np
import matplotlib.pyplot as plt
irisData = load_iris()
# Create feature and target arrays
X = irisData.data
y = irisData.target
# Split into training and test set
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size = 0.4, random_state=52)
neighbors = np.arange(1, 12)
train_accuracy = np.empty(len(neighbors))
test_accuracy = np.empty(len(neighbors))
# Loop over K values
for i, k in enumerate(neighbors):
    # print(i + " " + k)
    knn = KNeighborsClassifier(n_neighbors=k)
    knn.fit(X_train, y_train)
    # Compute training and test data accuracy
    train_accuracy[i] = knn.score(X_train, y_train)
    test_accuracy[i] = knn.score(X_test, y_test)
# Generate plot
plt.plot(neighbors, test_accuracy, label = 'Testing dataset Accuracy')
plt.plot(neighbors, train_accuracy, label = 'Training dataset Accuracy')

plt.legend()
plt.xlabel('n_neighbors')
plt.ylabel('Accuracy')
plt.show()
```

Output:-



(Assignment-08) Naïve Base Classifier

```
import numpy as np
import pandas as pd

# load iris dataset

from sklearn.datasets import load_iris
iris = load_iris()

# store the feature matrix (X) and response vector (y)
X = iris.data
y = iris.target

# splitting X and y into training and testing sets
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.4, random_state=1)

# training the model on training set
from sklearn.naive_bayes import GaussianNB
gnb = GaussianNB()
gnb.fit(X_train, y_train)

# making predictions on the testing set
y_pred = gnb.predict(X_test)

# comparing actual response values (y_test) with predicted response values (y_pred)
from sklearn import metrics

print("Gaussian Naive Bayes model accuracy(in %):", metrics.accuracy_score(y_test, y_pred)*100)
```

output :-

```
Gaussian Naive Bayes model accuracy(in %): 95.0
```


(Assignment-09) K-MeansClustering

```
import matplotlib.pyplot as plt

x = [4, 5, 10, 4, 3, 11, 14, 6, 10, 12]
y = [21, 19, 24, 17, 16, 25, 24, 22, 21, 21]

plt.scatter(x, y)
plt.show()
```

Output:-




```
from sklearn.cluster import KMeans

data = list(zip(x, y))
inertias = []

for i in range(1,11):
    kmeans = KMeans(n_clusters=i)
    kmeans.fit(data)
    inertias.append(kmeans.inertia_)

plt.plot(range(1,11), inertias, marker='o')
plt.title('Elbow method')
plt.xlabel('Number of clusters')
plt.ylabel('Inertia')
plt.show()
```


Output:-



```
kmeans = KMeans(n_clusters=2)
kmeans.fit(data)

plt.scatter(x, y, c=kmeans.labels_)
plt.show()
```

Output:-



output:-

```
# print data(feature) shape
```

```
# print the cancer labels (0:malignant, 1:benign)
```

[illegible]

```

# Import train_test_split function
from sklearn.model_selection import train_test_split

# Split dataset into training set and test set
X_train, X_test, y_train, y_test = train_test_split(cancer.data,
    cancer.target, test_size=0.3, random_state=109) # 70% training a
nd 30% test
#Import svm model
from sklearn import svm

#Create a svm Classifier
clf = svm.SVC(kernel='linear') # Linear Kernel

#Train the model using the training sets
clf.fit(X_train, y_train)

#Predict the response for test dataset
y_pred = clf.predict(X_test)
#Import scikit-learn metrics module for accuracy calculation
from sklearn import metrics

# Model Accuracy: how often is the classifier correct?
print("Accuracy:",metrics.accuracy_score(y_test, y_pred))
# Model Precision: what percentage of positive tuples are label
d as such?
print("Precision:",metrics.precision_score(y_test, y_pred))

# Model Recall: what percentage of positive tuples are labelled
as such?
print("Recall:",metrics.recall_score(y_test, y_pred))
Output:-

```

Accuracy: 0.9649122807017544

Precision: 0.9811320754716981

Recall: 0.9629629629629629
